TAKE-HOME FINDINGS · PRODUCT & ENGINEERING LEADERSHIP

Transcript Intelligence

100 calls, Feb–Apr 2026: how one known risk became a quarter-long outage story, and what each team should do next

Uma Maheswara G · AI · October 2026

WHEN IT HAPPENED

Warned five weeks early, broke on 10 March, still in conversations 49 days later

January

Ingestion pipeline fails. Same root cause. Disclosed to a customer on 7 Feb.

5 Feb

Sprint planning: “single processing node with no fallback.” Fix committed.

16 Feb

Standup: refactor 60% done, end-of-March deadline at 60/40 odds.

2 Mar

Sprint planning: the fix is cut back to a design spike.

10 Mar

Detect outage. Node hits a memory ceiling; ~6 hours of lost visibility.

April

The fix causes 4 new support cases. Last outage call 28 Apr.

In between, Comply v2 launched on 30 Mar and became the best week of the quarter, which hid how long the outage damage lasted.

02 · Dates from call metadata; quotes from internal call transcripts (notebook 04 §1a)

EXECUTIVE SUMMARY

The quarter in six answers

Question	Answer from the 100 calls
What was the issue?	Detect’s event ingestion ran on one node with no fallback; it failed on 10 March
What went wrong?	A known risk slipped for 5 weeks; customers weren’t notified; the fix broke other products
When?	Warnings from 5 Feb; outage 10–18 Mar; second wave in April; tail to 28 Apr
How was it handled?	46 calls, 21 of 32 accounts; only 2 of 30 support calls fully resolved
What was the fix?	Redundant nodes and a circuit breaker, tested only on paper so far
What next?	Fix communication first; review Protect before it fails the same way

WHAT WAS THE ISSUE

A known single point of failure, discussed in five calls and never fixed in time

“The event ingestion pipeline… it’s all running through a single processing node with no fallback. If that node goes down or gets overwhelmed, we lose event visibility entirely.”

Engineer, Detect sprint planning, 5 Feb · five weeks before the outage

6th

Where it ranked by mention count. The warning wasn’t loud.

#1, 8×

Where an AI risk read of internal calls ranks it, using only pre-outage data

WHAT WENT WRONG IN EACH DEPARTMENT

Every team saw part of it; no one saw all of it

Team	What went wrong	Evidence in the calls
Engineering	Known risk lost to feature work; the fix then broke other products	Fix cut to a design spike on 2 Mar; 4 April cases caused by remediation
Support	Had nothing to offer during a platform-side failure	5 outage-week calls ended unresolved and never improved
Account management	Absorbed escalations; some accounts went quiet with no follow-up	6 accounts with a bad support call and no AM call since
Product & comms	Customers learned of the outage late, and asked for RCAs and SLAs	Incident communication: customer tone −0.52, among the worst topics
Leadership	Signals sat in separate meetings with no shared view	Warnings spread across sprint, standup and renewal calls

HOW CUSTOMERS AND SUPPORT HANDLED IT

Staff stayed calm and polite, which hid how upset customers were

−0.37

Customer tone on support calls, while staff sat at +0.01

2 of 30

Support calls fully resolved; 23 ended with a workaround or escalation

+0.36

How much account calls improve from start to end

6

Accounts left with no follow-up after a bad support call

Account managers repair relationships during the call; support can’t fix a platform failure. Northstar and Ridgeline went about six weeks without contact.

HOW FAR IT REACHED

A six-hour outage that stayed in conversations for 49 days

46

of 100 calls discuss it: 17 account, 16 internal, 13 support

21 of 32

customer accounts touched by it

2.5×

its reach versus what call titles show (37 calls found by content, 15 by title)

Measure an incident by how long it stays in customer conversations, not by its downtime.

07 · Outage-related calls flagged by issue content, notebook 04 §2

WHAT WAS THE FIX

The fix worked for Detect, then caused a second wave of tickets

Remediation added redundant ingestion nodes and a circuit breaker. The failover has only been checked in a tabletop exercise, not a live node-failure test.

April case	What the fix broke	Account tone
Stratos Cloud	Circuit-breaker limit blocks large SIEM pushes	4.7 → 2.4
Summit Trust	Node rebalancing delays Detect alerts 20–45 min	3.8 → 2.4
Crestline Wealth	Shared-cluster change slows policy sync	3.8 → 2.4
Silverline Brands	Shared-cluster change pushes backups past their window	2.4 → 3.9

WHAT SENTIMENT SHOWS

The April “recovery” is an average of a launch success and a lingering outage



All three call types fell to about 2.2 in outage week. Comply v2’s launch week (30 Mar) was the best of the quarter, at about 4.8.

WHAT CUSTOMERS ASKED FOR

At-risk customers want honesty and guarantees, not new features

Top asks from at-risk accounts

Request	At-risk accts
Formal post-incident RCA reports	7
Proactive outage notifications	7
Enforceable uptime SLAs	6
Proof of pipeline resilience	6

Who is circling

SentinelShield appears in about **9%** of calls before the outage, **47%** during and **22%** after. Its pitch: native SIEM integrations and a 99.9% uptime SLA.

VaultEdge targets Protect on price; CyberNova was named in four Q1 lost deals.

The biggest overall request, multi-framework compliance reporting (21 accounts), already shipped in Comply v2, which explains the launch's lift.

THE WARNING FOR NEXT QUARTER

Protect today looks like Detect did in February

Critical: the backup scheduler can’t handle concurrent jobs at scale; queues reach 47+ minutes and Silverline’s backups already run into business hours.

Not scheduled: “single points of failure like Detect” in Protect, raised in three internal calls, discussed only.

Unproven: Detect’s new failover has passed only a tabletop exercise.

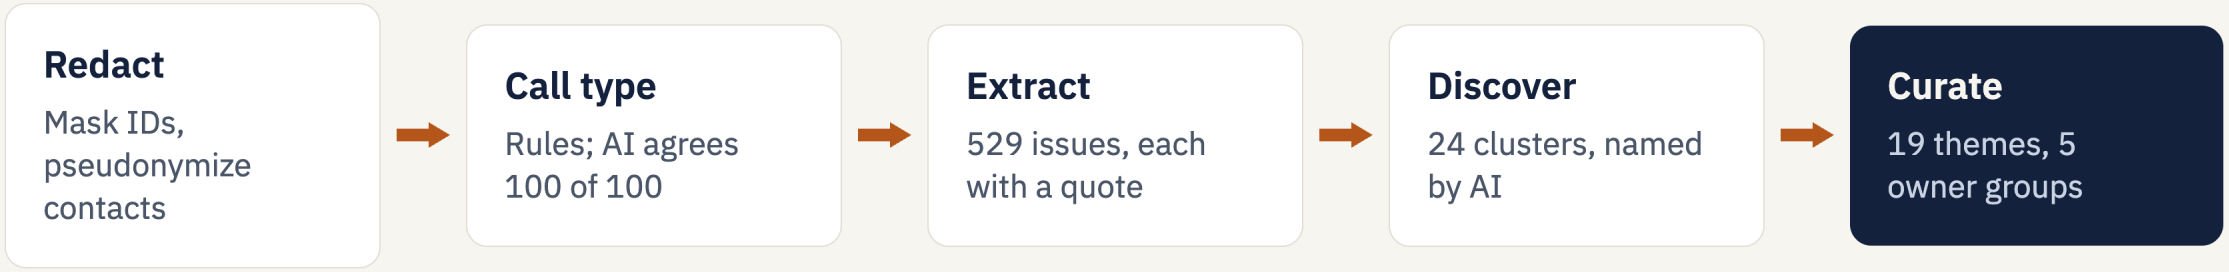
DIRECTION FOR EACH TEAM

One clear move per team, each backed by the calls

Team	Direction
Engineering	Run a live Detect failover test; fund a Protect single-point-of-failure review now
Release / QA	Treat remediation as a release: cross-product regression tests and change notes
Support	During platform incidents, hand callers to an incident lead; flag calls that end negative
Account mgmt	Close the 6 open loops this week with a call, not a ticket update
Product	Ship customer RCAs and proactive outage notices before the next feature
Sales	SentinelShield battlecard pairing the architecture fix with an SLA offer
Leadership	Review the cross-call risk radar weekly; track sentiment per topic

HOW THE ANALYSIS WAS DONE

Rules where they work, AI for meaning, a human for the categories



Check	Result
Evidence quotes found verbatim in transcripts	987 of 987
Sentiment conclusions under an independent model	All hold in direction (call-level r = 0.95)
Topic labels vs 25 hand-labeled calls	Sample drawn; labeling not yet done

THE CATEGORIES

The two outage themes lead, and appear in all three call types

Theme	Calls	Support	Account	Internal
Outage root cause & remediation	35	6	13	16
Outage business impact & trust	35	13	13	9
Compliance framework coverage	28	4	17	7
Incident communication	21	10	4	7
Competitive pressure	21	4	11	6
Onboarding & adoption	19	0	17	2
Renewal & pricing	16	0	14	2

Top 7 of 19 themes. A call can carry up to 4 themes; counted at confidence ≥ 0.6. Example quotes per theme are in notebook 02 §7.

HOW FAR TO TRUST THIS

What this analysis can and can’t claim

Can claim

- The outage timeline and its reach
- Every quote is real and checked
- Sentiment gaps hold under a second model

Can’t claim yet

- That the radar is reliable: one backtest
- Churn: no renewal outcomes in the data
- Statistical significance at 100 calls

The data is synthetic and transcripts are excerpts (median 29% of the meeting), so text is the source of truth and pace, accuracy and weekday aren’t used.